

# Applied Machine Learning

Regularization

**Reihaneh Rabbany**



# Learning objectives

Basic idea of

- overfitting and underfitting
- regularization (L1 & L2)
- MLE vs MAP estimation
- bias and variance trade off
- evaluation metrics & cross validation

# Previously...

Linear regression and logistic regression  
is linear too simple? what if it's not a good fit?  
how to increase the models expressiveness?

- create new nonlinear features
- is there a downside?

# Recall: Linear regression

model

$$f_w(x) = \underset{\in \mathbb{R}}{\hat{y}^{(n)}} = \underset{1 \times D}{w^T} \underset{D \times 1}{x^{(n)}}$$

cost

$$J(w) = \frac{1}{2} \sum_n \left( y^{(n)} - w^T x^{(n)} \right)^2$$

linear least squares (LLS)

Optimization

$$\sum_n (y^{(n)} - w^T x^{(n)}) x_d^{(n)} = 0 \quad \forall d$$

matrix notation

$$\underset{N \times 1}{\hat{y}} = \underset{N \times D}{X} \underset{D \times 1}{w}$$

$$J(w) = \frac{1}{2} \|y - Xw\|^2 \\ = \frac{1}{2} (y - Xw)^T (y - Xw)$$

$$X^T (y - Xw) = \vec{0}$$

$$\underset{D \times 1}{w^*} = \underset{D \times N}{(X^T X)^{-1}} \underset{N \times D}{X^T} \underset{N \times 1}{y}$$

# Recall: Linear regression

$$w^* = (X^T X)^{-1} X^T y$$

$D \times 1$     $D \times N$   $N \times D$     $N \times 1$

what if  $X^T X$  is **not invertible**?

add a small value to the diagonals, a.k.a. regularize

what if **linear fit is not the best**?

use nonlinear basis

# Recall: nonlinear basis functions

what if **linear fit is not the best**?



$$X = \begin{bmatrix} x_1^{(1)} & x_2^{(1)} & \cdots & x_D^{(1)} \\ \vdots & \vdots & \ddots & \vdots \\ x_1^{(N)} & x_2^{(N)} & \cdots & x_D^{(N)} \end{bmatrix}$$

replacing  $X$  with  $\Phi$

$$\Phi = \begin{bmatrix} \phi_1(x^{(1)}) & \phi_2(x^{(1)}) & \cdots & \phi_D(x^{(1)}) \\ \phi_1(x^{(2)}) & \phi_2(x^{(2)}) & \cdots & \phi_D(x^{(2)}) \\ \vdots & \vdots & \ddots & \vdots \\ \phi_1(x^{(N)}) & \phi_2(x^{(N)}) & \cdots & \phi_D(x^{(N)}) \end{bmatrix}$$

a (nonlinear) feature

Example



# Recall: nonlinear basis functions

what if **linear fit is not the best**?

replace original features in  $f_w(x) = \sum_d w_d x_d$





with nonlinear bases  $f_w(x) = \sum_d w_d \phi_d(x)$

linear least squares (LLS) solution  $w^* = (\Phi^T \Phi)^{-1} \Phi^T y$

How to avoid overfitting then?

regularize

$$\Phi_k(x) = x^k$$



# Regularization serves many purposes

$$\underset{D \times 1}{w^*} = (\underset{D \times N}{X^T} \underset{N \times D}{X})^{-1} \underset{N \times 1}{X^T y}$$

what if  $X^T X$  is **not invertible**?

add a small value to the diagonals, a.k.a. **regularize**

what if **linear fit is not the best**?

use nonlinear basis

How to avoid **overfitting** then? **regularize**

what if **we want a sparse model**?

do feature selection and only keep important parameters with **regularizing**


► Today, more on overfitting and regularization



# Recall: nonlinear basis functions


examples

original input is scalar  $x \in \mathbb{R}$




polynomial bases

$$\phi_k(x) = x^k$$



Gaussian bases


$$\phi_k(x) = e^{-\frac{(x-\mu_k)^2}{s^2}}$$




Sigmoid bases

$$\phi_k(x) = \frac{1}{1+e^{-\frac{x-\mu_k}{s}}}$$

# Example: Gaussian bases



$$\phi_k(x) = e^{-\frac{(x-\mu_k)^2}{s^2}}$$



our fit to data using 10 Gaussian bases

prediction for a new instance


$$f(\mathbf{x}') = \phi(\mathbf{x}')^T (\Phi^T \Phi)^{-1} \Phi^T \mathbf{y}$$

new instance


$\mathbf{w}$  found using LLS

features evaluated for the new point

## Example: Gaussian bases



$$\phi_k(x) = e^{-\frac{(x-\mu_k)^2}{s^2}}$$



our fit to data using **10 Gaussian bases**

why not more?


```
1 #x: N
2 #y: N
3 plt.plot(x, y, 'b.')
4 phi = lambda x, mu: np.exp(-(x-mu)**2)
5 mu = np.linspace(0,10,10) #10 Gaussians bases
6 Phi = phi(x[:,None], mu[None,:]) #N x 10
7 w = np.linalg.lstsq(Phi, y)[0]
8 yh = np.dot(Phi,w)
9 plt.plot(x, yh, 'g-')
```

# Example: Gaussian bases



$$\phi_k(x) = e^{-\frac{(x-\mu_k)^2}{s^2}}$$

using 50 bases




```
1 #x: N
2 #y: N
3 plt.plot(x, y, 'b.')
4 phi = lambda x, mu: np.exp(-(x-mu)**2)
5 mu = np.linspace(0,10,50) #50 Gaussians bases
6 Phi = phi(x[:,None], mu[None,:]) #N x 10
7 w = np.linalg.lstsq(Phi, y)[0]
8 yh = np.dot(Phi,w)
9 plt.plot(x, yh, 'g-')
```

# Example: Gaussian bases



$$\phi_k(x) = e^{-\frac{(x-\mu_k)^2}{s^2}}$$

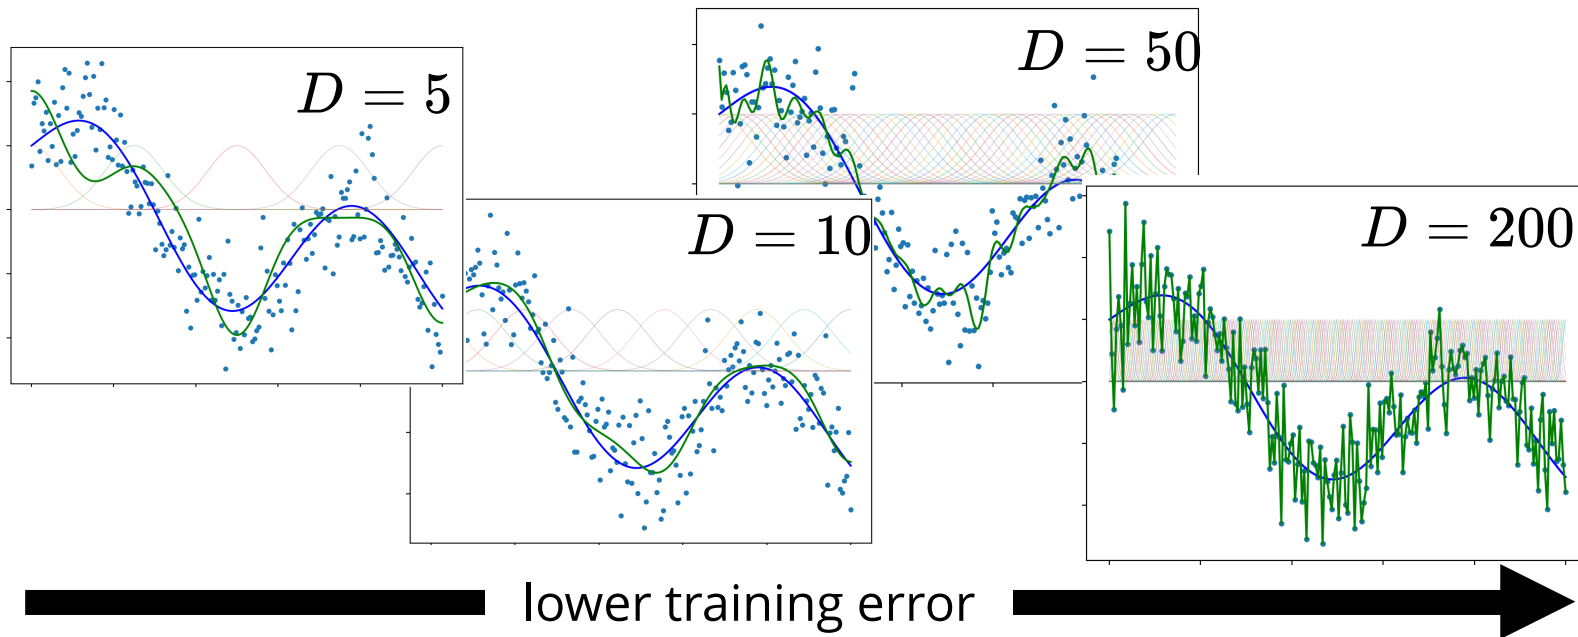
using 200, thinner bases ( $s=.1$ )



cost function  $J(w)$  is small and we have a "perfect" fit!

```
2 #y: N
3 plt.plot(x, y, 'b.')
4 phi = lambda x,mu: np.exp(-((x-mu)/.1)**2)
5 mu = np.linspace(0,10,200) #200 Gaussians bases
6 Phi = phi(x[:,None], mu[None,:]) #N x 10
7 w = np.linalg.lstsq(Phi, y)[0]
8 yh = np.dot(Phi,w)
9 plt.plot(x, yh, 'g-')
```


# Generalization



which one of these models performs better at **test time**?

# Overfitting

which one of these models performs better at **test time**?



# Model selection

how to pick the model with lowest expected loss / test error?

**regularization** bound the test error by bounding

- training error
- model complexity



USE a **validation set** (and a separate test set for final assessment)



use for model selection

use for final model assessment






# An observation

when overfitting, we often see large weights



dashed lines are  $w_d \phi_d(x) \quad \forall d$



**idea:** penalize large parameter values

# Ridge regression

L2 regularized linear least squares regression:

$$J(w) = \underbrace{\frac{1}{2} \|Xw - y\|_2^2}_{\substack{\text{sum of squared error} \\ \frac{1}{2} \sum_n (y^{(n)} - w^T x)^2}} + \underbrace{\frac{\lambda}{2} \|w\|_2^2}_{\substack{\text{(squared) L2 norm of } w \\ w^T w = \sum_d w^2}}$$

regularization parameter  $\lambda > 0$  controls the strength of regularization

a good practice is to not penalize the intercept  $\lambda(\|w\|_2^2 - w_0^2)$

# Ridge regression

we can set the derivative to zero  $J(w) = \frac{1}{2}(Xw - y)^T(Xw - y) + \frac{\lambda}{2}w^T w$

$$\frac{\partial Xw}{\partial w} = X^T \quad \text{Using matrix differentiation}$$
$$\frac{\partial w^T Xw}{\partial w} = 2Xw$$

$$\nabla J(w) = X^T(Xw - y) + \lambda w = 0$$

$$(X^T X + \lambda I)w = X^T y$$

when using gradient descent, this term reduces the weights at each step (**weight decay**)


$$w^* = (X^T X + \lambda I)^{-1} X^T y$$

the only part different due to regularization

$\lambda I$  makes it invertible!

we can have linearly dependent features (e.g.,  $D > N$ )

the solution will be unique!






# Example: polynomial bases


polynomial bases

$$\phi_k(x) = x^k$$

Without regularization:

- using  $D=10$  we can perfectly fit the data (high test error)






# Example: polynomial bases

polynomial bases


$$\phi_k(x) = x^k$$

with regularization:


- fixed  $D=10$ , changing the amount of regularization



$\lambda = 0$



$\lambda = .1$



$\lambda = 10$

# Data normalization

what if we scale the input features, using different factors  $\tilde{x}^{(n)} = \gamma_d x^{(n)} \forall d, n$

if we have **no regularization**:  $\tilde{w}_d = \frac{1}{\gamma_d} w_d \forall d$

everything remains the same because:  $\|Xw - y\|_2^2 = \|\tilde{X}\tilde{w} - y\|_2^2$

**with regularization**:  $\|\tilde{w}\|_2 \neq \|w\|_2^2$  so the optimal  $\mathbf{w}$  will be different!

features of different mean and variance will be penalized differently

normalization

$$\begin{cases} \mu_d = \frac{1}{N} x_d^{(n)} \\ \sigma_d^2 = \frac{1}{N-1} (x_d^{(n)} - \mu_d)^2 \end{cases}$$

makes sure all features have the same mean and variance  $x_d^{(n)} \leftarrow \frac{x_d^{(n)} - \mu_d}{\sigma_d}$

# Maximum likelihood

**previously:** linear regression & logistic regression maximize log-likelihood

linear regression

$$w^* = \arg \max_w p_w(y|x)$$

$$w^* = \arg \max_w p(y|x, w)$$

$$= \arg \max_w \prod_{n=1}^N \mathcal{N}(y; \Phi w, \sigma^2)$$

$$\equiv \arg \min \sum_n L_2(y^{(n)}, w^T \phi(x^{(n)}))$$

logistic regression

$$w^* = \arg \max_w p_w(y|x)$$

$$w^* = \arg \max_w p(y|x, w)$$

$$= \arg \max_w \prod_{n=1}^N \text{Bernoulli}(y; \sigma(\Phi w))$$

$$\equiv \arg \min \sum_n L_{CE}(y^{(n)}, \sigma(w^T \phi(x^n)))$$

idea: maximize the posterior instead of likelihood

$$p(w|y) = \frac{p(w)p(y|w)}{p(y)}$$

# Maximum a Posteriori (MAP)

use the Bayes rule and find the parameters with max posterior prob.

MAP estimate

$$p(w|y) = \frac{p(w)p(y|w)}{p(y)} \text{ the same for all choices of } w \text{ (ignore)}$$

$$w^* = \arg \max_w p(w)p(y|w)$$

$$\equiv \arg \max_w \log p(y|w) + \log p(w)$$

likelihood: original objective      prior

even better would be to estimate the posterior distribution  $p(w|y)$

- more on this later in the course!



# Gaussian prior

**Gaussian** likelihood and **Gaussian** prior

$$\mathcal{N}(\mu, \sigma) = \frac{1}{\sigma\sqrt{2\pi}} e^{-\frac{1}{2}\left(\frac{x-\mu}{\sigma}\right)^2}$$

$$w^* = \arg \max_w p(w)p(y|w) \equiv \arg \max_w \log p(y|w) + \log p(w)$$

$$\equiv \arg \max_w \log \mathcal{N}(y|w^T x, \sigma^2) + \sum_{d=1}^D \log \mathcal{N}(w_d, 0, \tau^2) \quad \text{assuming independent Gaussian (one per each weight)}$$

$$\equiv \arg \max_w \frac{-1}{2\sigma^2}(y - w^T x)^2 - \sum_{d=1}^D \frac{1}{2\tau^2} w_d^2 \equiv \arg \min_w \frac{1}{2}(y - w^T x)^2 + \sum_{d=1}^D \frac{\sigma^2}{2\tau^2} w_d^2$$

multiple data-points

$$\equiv \arg \min_w \frac{1}{2} \sum_n (y^{(n)} - w^T x^{(n)})^2 + \sum_{d=1}^D \frac{\lambda}{2} w_d^2 \quad \text{L2 regularization}$$

$\lambda = \frac{\sigma^2}{\tau^2}$   
↓

**L2- regularization** is assuming a **Gaussian prior** on weights

the same is true for logistic regression (or any other cost function)


# Laplace prior

another notable choice of prior is the Laplace distribution

minimizing negative log-likelihood  $\rightarrow -\sum_d \log p(w_d) = \sum_d \frac{1}{2\beta} |w_d| = \frac{1}{2\beta} ||w||_1$   
L1 norm of w

L1 regularization:  $J(w) \leftarrow J(w) + \lambda ||w||_1$  also called **lasso**

(least absolute shrinkage and selection operator)



$$p(w; \beta) = \frac{1}{2\beta} e^{-\frac{|w|}{\beta}}$$

notice the peak around zero

image: <https://stats.stackexchange.com/questions/177210/why-is-laplace-prior-producing-sparse-solutions>



# $L_1$ vs $L_2$ regularization

regularization path shows how  $\{w_d\}$  change as we change  $\lambda$

Lasso produces sparse weights (many are zero, rather than small)

red-line is the optimal  $\lambda$  from cross-validation

$$D = 8$$



$w_{d'}$

$w_d$

decreasing regularization coef.  $\lambda$   $\longrightarrow$


# $L_1$ vs $L_2$ regularization

$$\min_w J(w) + \lambda \|w\|_p^p$$

is equivalent to  $\min_w J(w)$  subject to  $\|w\|_p^p \leq \tilde{\lambda}$  for an appropriate choice of  $\tilde{\lambda}$


figures below show the constraint and the isocontours of  $J(w)$

optimal solution with L1-regularization is more likely to have zero components



# Subset selection

p-norms with  $p \geq 1$  are convex (easier to optimize)



penalizes the **number of** non-zero features


$$J(w) + \lambda \|w\|_0 = J(w) + \lambda \sum_d \mathbb{I}(w_d \neq 0)$$

a penalty of  $\lambda$  for each feature

performs feature selection

# Subset selection

p-norms with  $p \geq 1$  are convex (easier to optimize)



optimizing this is a difficult *combinatorial problem*:

- search over all  $2^D$  subsets

L1 regularization is a viable alternative to L0 regularization

# Bias-variance decomposition

for L2 loss

assume a true distribution  $p(x, y)$

the regression function is  $f(x) = \mathbb{E}_p[y|x]$

assume that a dataset  $\mathcal{D} = \{(x^{(n)}, y^{(n)})\}_n$  is sampled from  $p(x, y)$

let  $\hat{f}_{\mathcal{D}}$  be our model based on the dataset

what we care about is the **expected loss (aka risk)**

$$\mathbb{E}[(\hat{f}_{\mathcal{D}}(x) - y)^2]$$

all **blue** items are **random variables**

# Bias-variance decomposition

for L2 loss

what we care about is the **expected loss (aka risk)**

$$\mathbb{E}[(\hat{f}_{\mathcal{D}}(x) - y)^2] = \mathbb{E}[(\hat{f}_{\mathcal{D}}(x) - \mathbb{E}_{\mathcal{D}}[\hat{f}_{\mathcal{D}}(x)] - y + \mathbb{E}_{\mathcal{D}}[\hat{f}_{\mathcal{D}}(x)])^2]$$

$\hat{f}_{\mathcal{D}}(x)$        $y$

$\downarrow$                        $\downarrow$

$f(x) + \epsilon$

$\hat{f}_{\mathcal{D}}(x) + \mathbb{E}_{\mathcal{D}}[\hat{f}_{\mathcal{D}}(x)] - \mathbb{E}_{\mathcal{D}}[\hat{f}_{\mathcal{D}}(x)]$  add and subtract a term

the remaining terms evaluate to zero (check for yourself!)

$$= \underbrace{\mathbb{E}[(\hat{f}_{\mathcal{D}}(x) - \mathbb{E}_{\mathcal{D}}[\hat{f}_{\mathcal{D}}(x)])^2]}_{\text{variance}} + \underbrace{\mathbb{E}[(f(x) - \mathbb{E}_{\mathcal{D}}[\hat{f}_{\mathcal{D}}(x)])^2]}_{\text{bias}} + \underbrace{\mathbb{E}[\epsilon^2]}_{\text{unavoidable noise error}}$$



# Bias-variance decomposition

for L2 loss

the expected loss is decomposed to:

$$\mathbb{E}[(f(x) - \mathbb{E}_{\mathcal{D}}[\hat{f}_{\mathcal{D}}(x)])^2]$$


**bias:** how average over all datasets differs from the regression function

$$\mathbb{E}[(\hat{f}_{\mathcal{D}}(x) - \mathbb{E}_{\mathcal{D}}[\hat{f}_{\mathcal{D}}(x)])^2]$$

**variance:** how change of dataset affects the prediction

$$\mathbb{E}[\epsilon^2]$$

**noise error:** the error even if we used the true model  $f(x)$




different models vary in their trade off between error due to bias and variance



- simple models: often more biased
- complex models: often have more variance

image: P. Domingos' posted article

# Example: bias vs. variance



# Example: bias vs. variance





side note

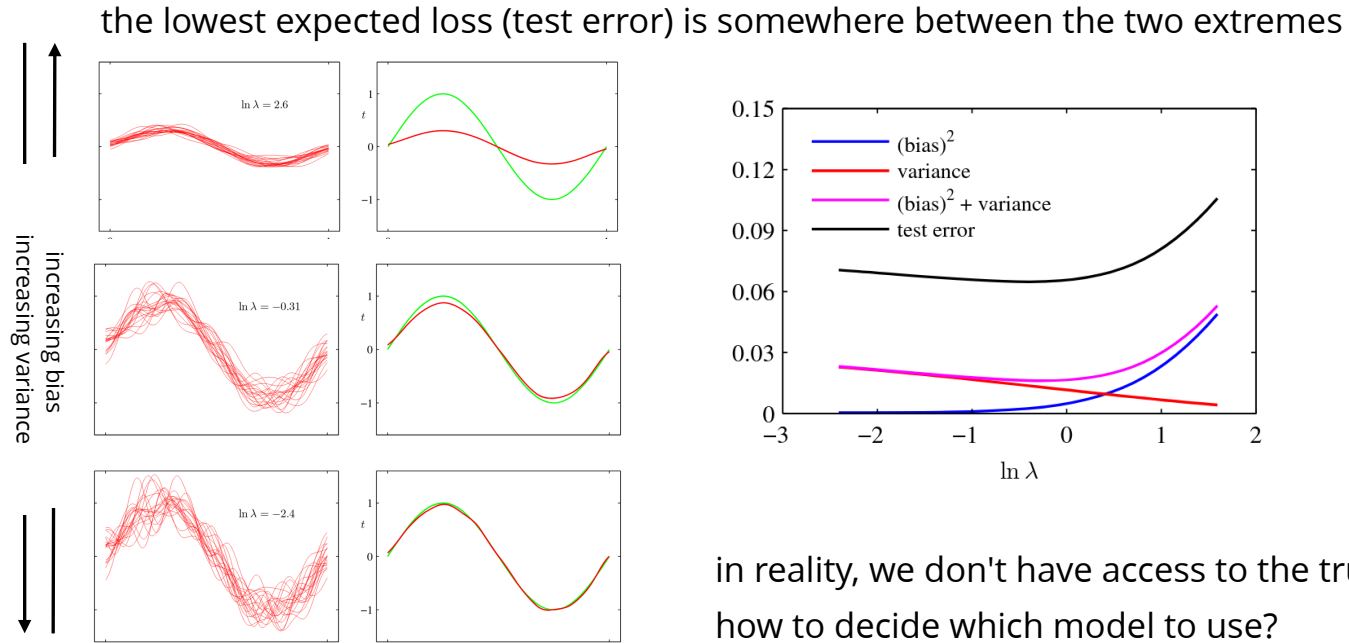
the average fit is very good, despite high variance

**model averaging:** uses "average" prediction of expressive models to prevent overfitting

using larger regularization penalty: higher bias - lower variance




# Example: bias vs. variance



# Big picture!

high variance in more complex models means that **test and training error can be very different**  
high bias in simplistic models means that **training error can be high**



# Model selection

how to pick the model with lowest expected loss / test error?

use a **validation set** (and a separate test set for final assessment)



use for model selection

use for final model assessment



**regularization** bound the test error by bounding

- training error
- model complexity


in the end we may have to use a validation set to find the right amount of regularization

# Cross validation

getting a more reliable estimate of test error using validation set

## K-fold cross validation(CV)

- randomly partition the data into  $K$  *folds*
- use  $K-1$  for training, and 1 for validation
- report average/std of the validation error over all folds




leave-one-out CV: extreme case of  $k=N$

# Cross validation

getting a more reliable estimate of test error using validation set

## K-fold cross validation(CV)

- randomly partition the data into  $k$  *folds*
- use  $k-1$  for training, and 1 for validation
- report average/std of the validation error over all folds



once the hyper-parameters are selected, we can use the whole set for training  
use test set for the **final** assessment

image credit: Thanh Nguyen et al'19



# Evaluation

**evaluation metric** can be different from the optimization objective

**confusion matrix** is a CxC table that compares truth-vs-prediction

for **binary classification**:

	Truth		Σ
	TP	FP	RP
Result	FN	TN	RN
Σ	P	N	

some **evaluation metrics**  
(based on the confusion table)

$$Accuracy = \frac{TP+TN}{P+N}$$


$$Error\ rate = \frac{FP+FN}{P+N}$$

$$Precision = \frac{TP}{RP}$$

$$Recall = \frac{TP}{P}$$


$$F_1\ score = 2 \frac{Precision \times Recall}{Precision + Recall}$$

type I vs type II error



# Evaluation

if we produce class score (probability)  
we can trade-off between type I & type II error




**goal:** evaluate class scores/probabilities (independent of choice of threshold)

Receiver Operating Characteristic **ROC curve**

**TPR** = TP/P (**recall**, sensitivity)

**FPR** = FP/N (**fallout**, false alarm)



# Summary

- complex models can have very different training and test error (*generalization gap*)
- regularization bounds this gap by penalizing model complexity
  - L1 & L2 regularization
  - probabilistic interpretation: different priors on weights
  - L1 produces sparse solutions (useful for feature selection)
- bias-variance trade off:
  - formalizes the relation between
    - training error (bias)
    - complexity (variance) and
    - and the test error (bias + variance)
  - not so elegant beyond L2 loss
- (cross) validation for model selection